



Ain Shams University

CSE489:

Selected Topics in Data Science

Arabic Voice Assistant Project

Presented to:

Dr. Alaa Hamdy

Prepared by

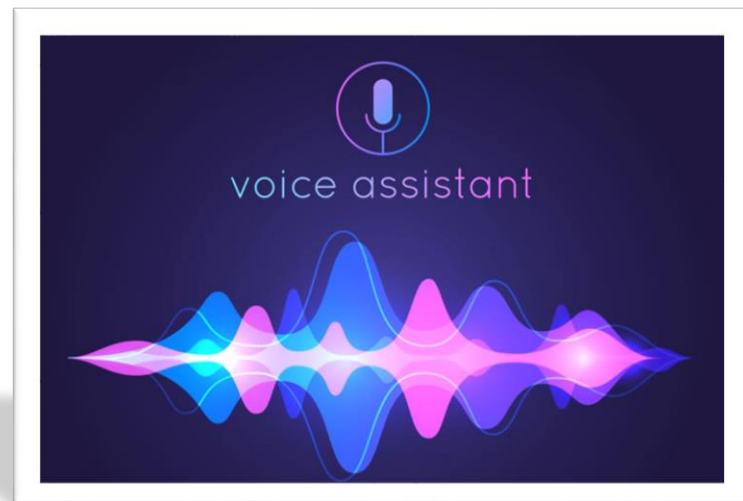
Team 9:

Adham Yasser

ID:20P9705

Laila Ihab

ID: 20P1005



Contents

1 Introduction	3
1.1 Project Overview.....	3
1.2 Objectives	3
1.3 Scope	3
1.4 Key Features	3
2 System Architecture	4
2.1 High-Level Architecture Diagram	4
2.2 Detailed Component Descriptions	5
3 Data and Model Training	6
3.1 Data Collection and Preprocessing.....	6
3.1.1 Trigger Word Dataset	6
3.1.2 Intent Classification Dataset.....	7
3.2 LSTM Model for Trigger Word Detection.....	7
3.2.1 Model Architecture	7
3.2.2 Training Process	8
3.2.3 Evaluation and Results	10
3.3 Arabic BERT Model for Intent Classification	12
3.3.1 Model Architecture	12
3.3.2 Training Process	13
3.3.3 Evaluation and Results	14
4 System Workflow.....	15
4.1 Trigger Word Detection	16
4.2 Speech-to-Text Conversion	16
4.3 Intent Classification.....	16
4.4 Response Generation with LLAMA 3.1	17
4.5 Text-to-Speech Conversion.....	17
5 Implementation Details	17
5.1 Software and Tools Used	17
5.2 API Integrations	19
5.2.1 Google Speech-to-Text	19

5.2.2 Google Text-to-Speech.....	19
5.3 Prompt Engineering with LangChain	19
6 Challenges and Solutions.....	20
6.1 Data Collection Challenges	20
6.2 Model Training Challenges	20
6.3 Integration Challenges.....	20
7 Future Work and Improvements	20
7.1 Potential Enhancements.....	20
7.2 Additional Features.....	21
8 Conclusion	21
8.1 Summary of Achievements	21
8.2 Final Thoughts.....	22
9 Appendices.....	22
9.1 Source Code.....	22
9.2 Colab Notebook.....	22

1 Introduction

1.1 Project Overview

The Arabic Smart Assistant project is an innovative application designed to understand and respond to Arabic spoken language. The system uses a combination of machine learning models and advanced natural language processing (NLP) techniques to detect a specific trigger word (whistling), transcribe spoken Arabic into text, classify the intent of the transcribed text, and generate appropriate responses in Arabic. The project integrates several technologies, including an LSTM model for trigger word detection, Google Speech-to-Text for transcription, an Arabic BERT model for intent classification, and a LLAMA large language model (LLM) for response generation. The final output is delivered to the user via Google Text-to-Speech, providing a seamless conversational experience.

1.2 Objectives

- Develop a robust system capable of detecting a specific trigger word (whistling) using a custom-trained LSTM model.
- Implement an accurate transcription of spoken Arabic to text using Google Speech-to-Text.
- Classify the transcribed Arabic text into predefined intents using a fine-tuned Arabic BERT model.
- Generate relevant and coherent Arabic responses using a LLAMA large language model, enhanced through prompt engineering and the LangChain framework.
- Deliver the generated response in spoken Arabic via Google Text-to-Speech, creating a user-friendly interaction.

1.3 Scope

The scope of this project encompasses the development and integration of various machine learning models and NLP techniques to create a fully functional Arabic smart assistant. This includes the creation of a custom dataset for trigger word detection, the fine-tuning of a pre-trained Arabic BERT model for intent classification, and the use of a LLAMA LLM for generating responses. The project also involves the implementation of speech-to-text and text-to-speech services to ensure smooth interaction with the user. The system is specifically designed to understand and respond to spoken Arabic, making it highly relevant for Arabic-speaking users.

1.4 Key Features

- **Custom Trigger Word Detection:** Utilizes a self-produced dataset along with the Common Voice dataset to train an LSTM model for detecting the trigger word, which is a whistle.

- **Arabic Transcription:** Employs Google Speech-to-Text to convert spoken Arabic into text.
- **Intent Classification:** Leverages a fine-tuned Arabic BERT model to classify the transcribed Arabic text into specific intents, enabling the assistant to understand user queries effectively.
- **Dynamic Response Generation:** Uses a LLAMA large language model, with the help of prompt engineering and the LangChain framework, to generate appropriate and contextually relevant Arabic responses.
- **Seamless User Interaction:** Converts the generated Arabic text into speech using Google Text-to-Speech, providing a natural and intuitive user experience.

2 System Architecture

2.1 High-Level Architecture Diagram

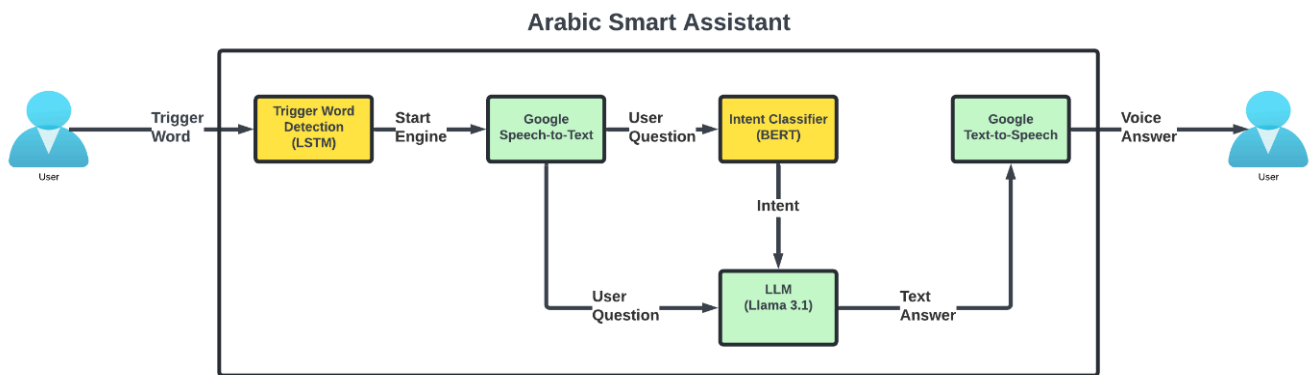


Figure 1: System Architecture

The high-level architecture of the Arabic Smart Assistant project outlines the interaction between various components to achieve the desired functionality. At its core, the system processes spoken Arabic input, detects trigger word, identifies the user's intent, generates an appropriate response, and then communicates that response back to the user. The main components are:

- **Speech Input:** The user communicates via voice, which is captured by the system.
- **Speech Recognition (Google Speech-to-Text):** The speech is converted into text using the Google Speech-to-Text API, specifically tailored to handle Arabic. This transcription is the starting point for subsequent processing.
- **Trigger Word Detection (LSTM Model):** After transcription, the system checks for predefined trigger words using a Long Short-Term Memory (LSTM) model.

- **Intent Classification (Arabic BERT Model):** Once a trigger word is detected, the input text is passed to a pre-trained Arabic BERT model that identifies the user's intent.
- **Response Generation (LLAMA LLM):** Based on the identified intent, a response is generated using the LLAMA Large Language Model (LLM).
- **Output:** The response can be delivered back to the user as synthesized speech.

This architecture diagram will show these components and how data flows between them, ensuring a smooth and responsive interaction process.

2.2 Detailed Component Descriptions

1. Speech Input (Google Speech-to-Text):

This component captures spoken input in Arabic from the user via a microphone. The audio stream is then processed by the Google Speech-to-Text API, which transcribes the spoken words into text with a focus on accuracy in the Arabic language. The choice of Google's API is based on its support for diverse dialects of Arabic and its ability to handle continuous speech in real time.

2. Trigger Word Detection (LSTM Model):

The transcribed text is fed into an LSTM model that monitors the input for predefined trigger words. A trigger word is essentially a specific phrase or keyword (such as "Hey Assistant") that activates the assistant's core functionality. The LSTM model is well-suited for this task as it excels in capturing time dependencies and sequential patterns in the input text, making it effective for recognizing trigger phrases within a continuous speech stream.

3. Intent Classification (Arabic BERT Model):

If the LSTM model detects a trigger word, the transcribed text is passed to an Arabic BERT model. The BERT model is fine-tuned to classify the user's intent based on the text input. This intent could range from setting a reminder, asking for the weather, playing music, or providing other services. BERT's transformer architecture allows it to understand the contextual meaning of words and phrases, which is essential for handling the nuances of Arabic syntax and grammar.

4. Response Generation (LLAMA LLM):

Once the intent is classified, the system relies on the LLAMA large language model to generate a natural-language response. LLAMA has been fine-tuned for Arabic, ensuring that its responses are coherent and contextually appropriate. The LLM processes the user's request and generates a corresponding answer, which could involve querying external APIs (e.g., weather services) or providing predefined responses.

5. Output:

The generated response is either converted back into speech using a text-to-speech system and displayed as text on the user's interface. The system uses Arabic text-to-speech synthesis to provide a smooth, human-like voice experience for the user.

3 Data and Model Training

3.1 Data Collection and Preprocessing

3.1.1 Trigger Word Dataset

For training the LSTM model to detect the trigger word, a specialized dataset of voice recordings is required. The dataset includes:

- **Voice Samples:** Collected from different speakers, they are recorded in various settings to introduce background noise, simulating real-world conditions.
- **Trigger word:** Specific word that should activate the system which is a whistle.
- **Non trigger words:** Voice samples that don't contain the trigger word to allow the model to differentiate between invalid and valid inputs.

The data collection and preprocessing for the trigger word dataset involve several key steps that ensure the dataset is both diverse and well-structured for model training.

1. **Audio Collection:** The first step is gathering raw audio recordings containing examples of the trigger word (wake word). These recordings are collected from various users in different environments, capturing variations in accents, noise levels, and speaking styles. This diversity in data helps ensure that the model can generalize well across different speakers and background conditions.
2. **Dataset Structuring:** Once the audio recordings are collected, they are organized into a structured format suitable for model training. Each audio file is annotated with metadata, including the file path, sample rate, and corresponding labels. This structure allows the training pipeline to efficiently access and process the audio data during the model's learning phase.
3. **Data Augmentation:** To enhance the model's ability to generalize, data augmentation techniques are applied to the collected audio samples. This process introduces variations such as pitch shifting, time stretching, and adding background noise to the recordings. By augmenting the dataset, the model is exposed to a wider range of audio conditions, making it more robust and capable of detecting trigger words in diverse real-world scenarios.
4. **Chunking Audio:** Long audio recordings are segmented into smaller, manageable chunks. This step ensures that the model trains on uniform-length

audio segments, improving its ability to focus on the relevant portions of the recordings where the trigger words appear. Training with smaller chunks also enhances the model's precision in detecting trigger words, reducing distractions from irrelevant audio segments.

5. **Dataset Splitting:** Finally, the dataset is split into training, validation, and testing sets. This ensures that the model is trained on one subset of the data, validated on another for tuning, and evaluated on a completely unseen set to measure its performance.

3.1.2 Intent Classification Dataset

The intent classification dataset undergoes a series of preprocessing steps to prepare the Arabic text data for use with a BERT model. Initially, the **ArabertPreprocessor** is employed to clean and normalize the Arabic text, removing unwanted characters and standardizing the format. This ensures that the model can learn from consistent input data. The **BERT tokenizer** and model are loaded, both pretrained on Arabic tweets ("aubmindlab/bert-base-arabertv02-twitter"), making them well-suited for this specific text domain. The preprocess function from the **ArabertPreprocessor** is applied to each entry in the 'text' column of the dataset to further standardize the text before tokenization. The dataset is then split into training and validation sets using **train_test_split** with stratification based on the 'intent' column to maintain the balance of classes in both sets. The 'intent' labels are mapped to numerical values to prepare them for classification. Both the training and validation texts are then tokenized using the BERT tokenizer, with padding and truncation applied to ensure a uniform sequence length across the dataset. This ensures that the text is in a format suitable for input into the TensorFlow-based BERT model. These steps are critical in preparing the dataset for effective training and evaluation, enabling the BERT model to learn meaningful patterns for intent classification in Arabic text.

3.2 LSTM Model for Trigger Word Detection

3.2.1 Model Architecture

The code provided constructs an LSTM (Long Short-Term Memory) model designed for detecting trigger words, specifically implemented in the **LSTMWakeWord** class. This model architecture leverages LSTM, a recurrent neural network (RNN) variant known for handling sequences of data, such as audio signals, while effectively mitigating the vanishing gradient problem, making it suitable for detecting patterns over long-time dependencies.

The **LSTMWakeWord** class inherits from `nn.Module` in PyTorch, allowing it to be integrated seamlessly into the PyTorch framework. Its core components include:

1. **Layer Normalization:**

The first step in the architecture is `nn.LayerNorm`, which normalizes the features of the input MFCC (Mel-frequency cepstral coefficients) to stabilize the training process and mitigate internal covariate shifts. This step ensures that the model converges faster by scaling the features before feeding them into the LSTM network.

2. **LSTM Layer:**

The LSTM layer is defined using `nn.LSTM`, where each input is processed through time to learn sequential dependencies. The key parameters for the LSTM layer include:

- a. `input_size`: The size of each input feature vector, which is set to the feature size (in this case, MFCC feature size = 40).
- b. `hidden_size`: The number of LSTM cells in each layer (defined in the training script). It controls how much memory the model holds.
- c. `num_layers`: The depth of the LSTM, determining how many layers of LSTM cells are stacked on top of each other. This adds complexity and allows the model to capture deeper patterns.
- d. `dropout`: Applied to prevent overfitting by randomly dropping units during training.
- e. `bidirectional`: Defines whether the LSTM should process the input sequence in both forward and backward directions. In this case, it's set to `False`, making it a unidirectional LSTM.

3. **Linear Classifier:**

After processing the input sequence through the LSTM layers, the final hidden state (`hn`) is passed through a linear classifier (`nn.Linear`). This layer reduces the dimensionality to the number of output classes, which is 1, as this is a binary classification task to detect whether a trigger word is present.

4. **Hidden State Initialization:**

The `_init_hidden` method initializes the hidden and cell states of the LSTM, which are required to maintain the sequence memory. The number of directions and layers are accounted for in the shape of the hidden states, ensuring the model can start learning patterns from scratch.

3.2.2 Training Process

The training process occurs in the `train.py` script. Several key components are used in this process:

1. **Data Preprocessing with MFCC and Augmentation:**

The `WakeWordData` class loads and processes the audio data. The MFCC features are extracted using the `MFCC` class, converting the raw audio waveform into a representation that reflects how humans perceive sound. In the training phase, the data undergoes augmentation (`SpecAugment`), adding random

transformations to the audio input. This technique increases the model's robustness to variations in the audio signal, making it more generalizable.

2. **Loss Function and Optimizer:**

During training, the model uses binary cross-entropy loss with logits (BCEWithLogitsLoss). This loss is appropriate for binary classification tasks, where the output needs to be interpreted as a probability between 0 and 1. The AdamW optimizer, a variant of the Adam optimizer, is used for updating the model parameters, and it includes weight decay to prevent overfitting.

3. **Learning Rate Scheduler:**

A learning rate scheduler (ReduceLROnPlateau) monitors the training performance and reduces the learning rate if the model's performance plateaus, preventing unnecessary fluctuations and ensuring smoother convergence.

4. **Evaluation:**

After each epoch, the model is evaluated on the test data. The test function calculates the binary accuracy by rounding the sigmoid outputs of the LSTM. A classification report is generated using metrics like precision, recall, and F1-score to provide insights into the model's performance.

5. **Checkpointing:**

The model checkpoints are saved based on test accuracy. This ensures that the best-performing model during training is preserved and can be reloaded later.

3.2.3 Evaluation and Results




Figure 2: LSTM Accuracy

The evaluation and results section centers on the model's training and testing process, where performance metrics such as accuracy are gathered and analyzed. The provided code is designed to train and evaluate a model for wake word detection using LSTMs. Let's break down the process, explaining why each component is essential.

The core of the process begins with the `train` function. During training, the model iterates over batches of input data (mfcc features) and their corresponding labels. These features represent the audio signal's MFCC (Mel-frequency Cepstral Coefficients), which are commonly used in speech and audio recognition tasks. Both the features and labels are moved to the appropriate device (either GPU or CPU) to leverage hardware acceleration when available.

1. Optimization and Loss Calculation:

The optimizer's gradients are reset using `optimizer.zero_grad()`, followed by a forward pass through the model to predict outputs (`output = model(mfcc)`). The loss between the predicted and actual labels is computed using `BCEWithLogitsLoss`, which is suited for binary classification tasks like wake word detection. The backpropagation step (`loss.backward()`) ensures that the model learns by adjusting its parameters based on the computed gradients.

2. **Prediction and Accuracy:**

After calculating the loss, predictions are generated by applying the sigmoid function to the model's outputs, converting the raw logits into probabilities. These predictions are compared with actual labels to compute accuracy, using a helper function `binary_accuracy`. The binary rounding of predictions ensures that values above 0.5 are classified as positive, while others are negative. Additionally, the classification report is generated to give detailed metrics like precision, recall, and F1-score.

3. **Epoch-wise Training and Testing:**

The main function handles the overall process, beginning by setting up the datasets for training and testing. `train_loader` and `test_loader` are `DataLoader` instances that batch the data and handle other optimizations for faster processing. An LSTM model (`LSTMWakeword`) is instantiated with parameters like the number of layers, hidden size, and dropout rate. The model can be initialized from a pre-trained checkpoint if specified.

During each epoch, the model is trained (`train`) and evaluated (`test`). The training and testing accuracy is tracked over time, and the best results are saved for checkpointing. The learning rate is dynamically adjusted by `ReduceLROnPlateau`, which reduces the learning rate when performance plateaus. This helps the model avoid getting stuck in local minima during optimization.

4. **Saving the Best Model:**

The code includes logic to save the model's state when it achieves a new best test accuracy, ensuring that the most optimal version of the model is preserved. This allows for later retrieval and deployment of the best-performing model.

3.3 Arabic BERT Model for Intent Classification

3.3.1 Model Architecture

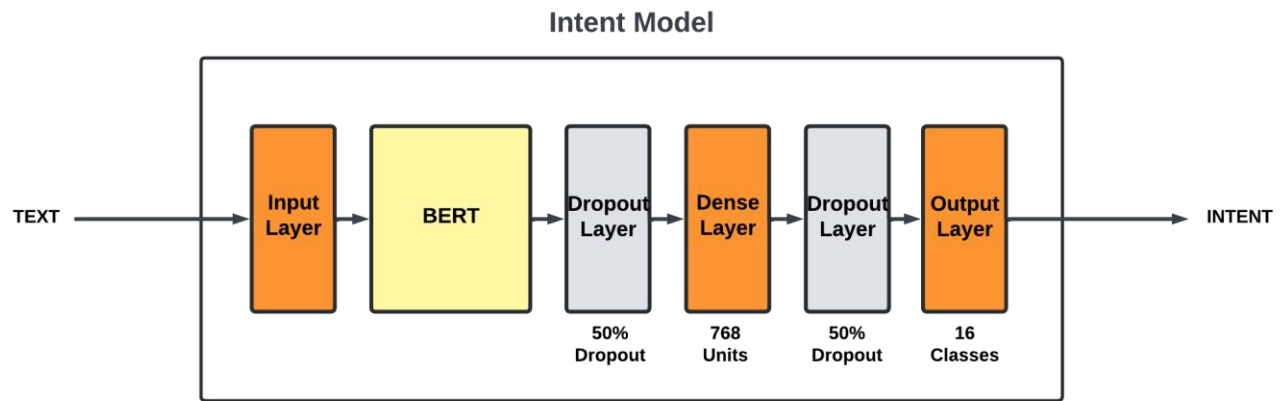


Figure 3: Intent Model Architecture

The **Arabic BERT Model for Intent Classification** leverages the pre-trained "aubmindlab/bert-base-arabertv02-twitter" model to perform classification of Arabic text into different intents. The architecture starts by tokenizing the input text using a BERT tokenizer specifically designed for Arabic, which transforms the text into a format suitable for processing by the BERT model. The core of the architecture is a fine-tuned BERT model that processes the input text and generates a meaningful representation of the content. This representation is passed through multiple layers, including a dropout layer for regularization and a dense layer with a tanh activation function to extract high-level features. The model ends with a softmax output layer, which classifies the input text into one of the predefined intent classes. The use of dropout ensures that the model generalizes well, preventing overfitting and improving performance on unseen data.

Explanation of Each Part:

1. **Tokenizer:** The BERT tokenizer breaks the Arabic text into tokens that the BERT model can understand. Without tokenization, the model would not be able to interpret the text.
2. **BERT Model:** The pre-trained BERT model captures the meaning of the Arabic text by generating high-quality embeddings, which are useful for downstream tasks like intent classification.
3. **Dropout Layer:** Dropout is used here to prevent overfitting by randomly disabling a fraction of neurons during training, making the model more generalizable.
4. **Dense Layer:** This fully connected layer, with a tanh activation function, processes the BERT output to better separate and classify the input data.

5. **Second Dropout Layer:** The second dropout layer further helps in regularization by introducing randomness in the model, ensuring that no particular set of features dominates.
6. **Output Layer:** The softmax activation in the output layer transforms the dense layer's output into a probability distribution over the possible intent classes, allowing the model to predict the most likely intent for a given text.

By combining these components, the model is fine-tuned to effectively classify Arabic text into specific user intents, making it a crucial part of the smart assistant's ability to understand and respond appropriately.

3.3.2 Training Process

The **training process** of the Arabic BERT model for intent classification is designed to optimize the model for accurate classification of Arabic text into predefined intents. The model uses the **Sparse Categorical Crossentropy** loss function, which is suitable for multi-class classification tasks where labels are integer-encoded. The accuracy during training is measured using the **Sparse Categorical Accuracy** metric, which calculates how often the predicted class matches the actual class. The training process is driven by the **Adam optimizer**, which is known for its efficiency and adaptability, using a small learning rate of $5e-6$ to ensure stable convergence.

During training, the model is fed with tokenized input text and their corresponding labels. The training process spans across **100 epochs**, ensuring that the model has ample opportunity to learn from the data. The use of validation data helps monitor the model's performance on unseen data during training, allowing for adjustments to prevent overfitting. The data is shuffled after each epoch to introduce randomness into the training process, which enhances the model's generalization ability. This training setup ensures that the model learns to accurately predict user intents from Arabic text.

Explanation of Each Part:

1. **Loss Function:** Sparse Categorical Crossentropy is chosen because the classification problem involves predicting one of several intent classes, and the labels are integer-encoded.
2. **Accuracy Metric:** Sparse Categorical Accuracy is used to track the proportion of correct predictions, providing insight into the model's performance during training and validation.
3. **Adam Optimizer:** Adam is an advanced optimization algorithm that adjusts the learning rate dynamically. Its use here ensures that the model can converge efficiently, even with a small learning rate, by optimizing the weights for accurate predictions.
4. **Training Process:** The fit method facilitates the training process by iterating over the training data for 100 epochs, continuously updating the model weights based

on the loss function. The inclusion of validation data ensures that the model does not overfit, and the use of shuffling adds randomness to improve generalization.

This process ensures that the model becomes highly capable of classifying intents from Arabic text by learning from both training data and validation feedback.

3.3.3 Evaluation and Results

The performance of the Arabic BERT model for intent classification is visualized using two key metrics: **accuracy** and **loss** over 100 training epochs.

1. Training Accuracy:

- The green line represents the **training accuracy**, showing that the model's accuracy improves steadily as the epochs progress, achieving nearly 100% after about 70 epochs.
- The blue line represents the **validation accuracy**, which follows a similar upward trend, although with slight fluctuations. The validation accuracy stabilizes after about 50 epochs and achieves around 90% towards the final epochs.




Figure 4: Intent Model Accuracy

- This suggests that the model generalizes well on the validation set and effectively learns the intended task, with minimal overfitting since both training and validation accuracy converge closely.

2. Training Loss:

- The training loss (green line) decreases sharply during the initial epochs and continues to decrease at a slower rate over time. By the 100th epoch, the loss is close to zero.
- The validation loss (blue line) also decreases steadily, indicating that the model is effectively minimizing the difference between predicted and actual labels.
- However, there are some fluctuations in validation loss, suggesting slight variations in performance across different batches or noise in the validation data, but overall, the validation loss follows a decreasing trend similar to the training loss.

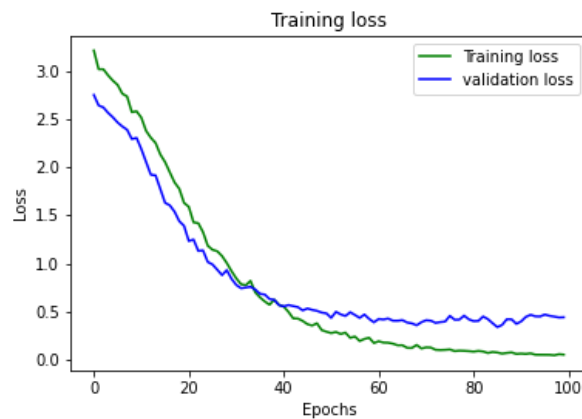


Figure 5: Intent Model Loss

These results reflect a well-performing model with a high degree of accuracy and low loss, indicating successful training and validation of the Arabic BERT model for intent classification. The model is learning effectively without major signs of overfitting, as evidenced by the close alignment between training and validation metrics.

4 System Workflow

The Arabic Smart Assistant follows a structured workflow comprising several key components to deliver a seamless interaction experience from detecting a trigger word to generating and delivering a spoken response. Below is a detailed description of each step in the system's workflow:

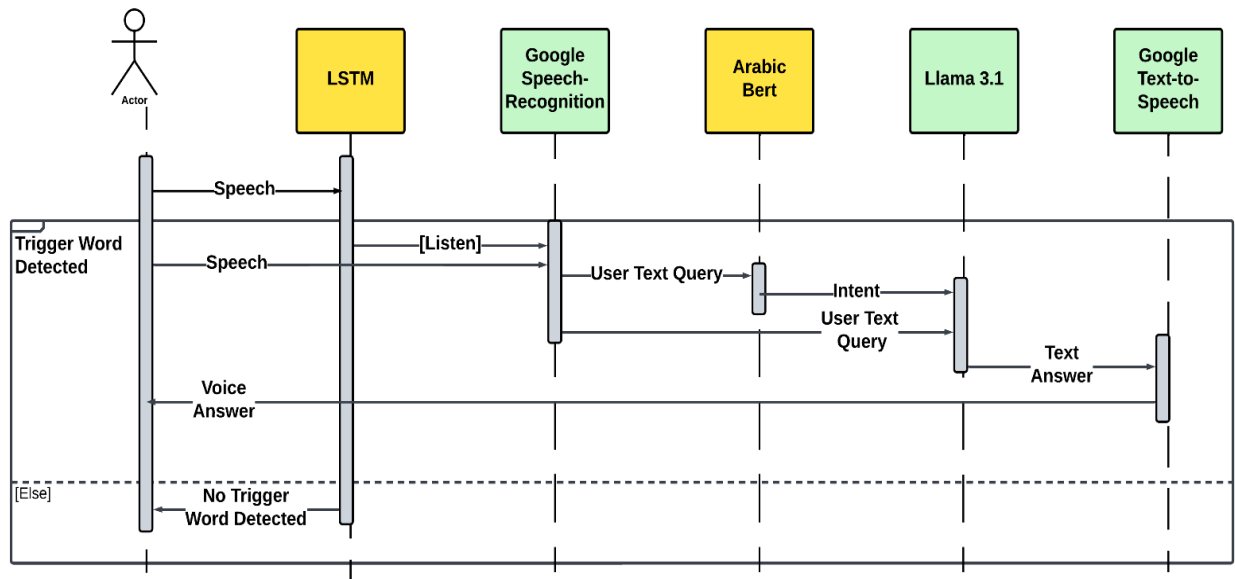


Figure 6: System Sequence Diagram

4.1 Trigger Word Detection

The system first listens for a specific trigger word, which in this case is a whistle sound. An LSTM model is trained to detect this sound, using both a custom dataset and the Common Voice dataset for greater variability. The LSTM model processes audio input in real-time and predicts whether the whistle trigger word is present. If detected, the system proceeds to the next step. This detection ensures the assistant only activates when prompted, reducing unnecessary processing and enhancing system efficiency.

Why it's used: LSTMs are excellent at capturing temporal dependencies in sequential data, making them well-suited for recognizing trigger words from audio signals over time.

4.2 Speech-to-Text Conversion

Once the trigger word (whistling) is detected, the system transitions to capturing and transcribing the user's spoken input. The speech-to-text conversion is performed using the **speech_recognition** library, which provides an interface to various speech recognition engines, including Google Speech-to-Text. This process converts the Arabic spoken words into text format, allowing the system to further process and understand the user's request. The accuracy of this conversion is critical, as it directly impacts the subsequent intent classification and response generation steps.

4.3 Intent Classification

The Arabic transcribed text is then passed to an Arabic BERT model that has been trained to classify the user's intent. The model is trained on an Arabic dataset from Kaggle, where it learns

to categorize text into different intent labels, such as asking for weather updates, sending messages, or retrieving information.

Why it's used: BERT models are highly effective at understanding context in text and have been fine-tuned for specific NLP tasks, such as intent classification. Arabic BERT is specifically tailored for handling Arabic text, making it ideal for understanding user commands.

4.4 Response Generation with LLAMA 3.1

Once the intent is determined, the input text and its classified intent are passed into the LLAMA language model (LLM) for generating an appropriate response. Using prompt engineering and the Langchain framework, the system structures the input for the LLAMA model to produce a well-formed answer. The response is dynamically generated based on the user's input and the identified intent.

Why it's used: LLAMA, as a large language model, has a powerful ability to generate human-like responses and adapt to different inputs. Using prompt engineering with Langchain allows for precise control over how the model generates answers, making the assistant versatile in responding to various queries.

4.5 Text-to-Speech Conversion

After generating the appropriate response based on the user's intent, the system converts the text-based response back into speech to provide an audible output. This is achieved using the **gtts** (Google Text-to-Speech) library, which allows the conversion of text into spoken words in Arabic. This last step ensures that the user receives a clear and understandable spoken response, completing the interaction cycle.

5 Implementation Details

5.1 Software and Tools Used

This section outlines the software and tools utilized in building the system, specifically for **trigger-word detection**, **intent classification**, and **language model (LLM) integration**. Below is a detailed description of each tool and its role in the system:

- **PyTorch:**
 - PyTorch is used to build the **LSTM-based wakeword detection model** (LSTMWakeWord). This model processes audio data to detect a specific trigger word (e.g., whistling).
 - The LSTM model is designed with multiple layers and an option for **bidirectional processing**, which allows it to capture patterns from both past and future audio sequences.

- A **LayerNorm** operation is applied to the input data to ensure stability during training by normalizing the feature values.
- The model outputs its prediction through a **Linear Classifier** that determines whether the input audio contains the wakeword.
- **TensorFlow:**
 - TensorFlow is used for fine-tuning a **pre-trained BERT model** for **intent classification** (IntentModel class).
 - This model is loaded on a **GPU** for faster inference, allowing the system to quickly classify user intents based on the input text.
 - The **BERT Tokenizer** is employed to convert input text into a format that the BERT model can process. The input is limited to a **maximum length** of 32 tokens to ensure efficient computation.
 - The BERT model is pre-trained on Arabic Twitter data and fine-tuned on specific intents such as 'Question,' 'Send Emails,' 'Open App,' etc.
- **LangChain:**
 - LangChain is utilized to integrate a **large language model (LLM)** to handle more complex interactions. The llm class uses **ChatGroq** with the **LLAMA-3.1-70B model** to generate dynamic and context-aware responses.
 - The system maintains **conversation history** using the **ChatMessageHistory** feature, ensuring continuity in interactions.
 - **PromptTemplates** are used to structure user inputs and system responses, guiding the LLM to generate relevant and coherent output.
- **Sonopy (MFCC):**
 - **Sonopy** is employed for extracting **Mel-Frequency Cepstral Coefficients (MFCCs)**, which are vital audio features for the LSTM-based wakeword detection model.
 - MFCCs capture the frequency characteristics of audio signals and are commonly used in speech and sound recognition tasks, providing critical input data for the wakeword model.

5.2 API Integrations

5.2.1 Google Speech-to-Text

The project integrates Google Speech-to-Text through the **speech_recognition** library, enabling the conversion of Arabic spoken language into text. This API is well-suited for handling various Arabic dialects and accents, ensuring accurate transcription of user inputs. The integration process involves configuring the API keys and setting up the library to capture audio input, which is then processed to produce text output.

5.2.2 Google Text-to-Speech

For generating the spoken response, the project utilizes the **gtts** library, which acts as a wrapper for Google Text-to-Speech. This API converts the text generated by the LLAMA LLM into natural-sounding Arabic speech. The integration of this API is straightforward, requiring only the installation of the gtts library and a simple function call to convert text into speech, which is then played back to the user. This allows for the creation of dynamic, real-time verbal responses, enhancing the overall user experience.

5.3 Prompt Engineering with LangChain

The **Prompt Engineering with LangChain** for this project leverages structured prompt templates to guide the interaction between the Arabic smart assistant and the user. The class "IIm" integrates ChatGroq's powerful language model with specific examples of how to handle different user intents, such as reading calendar events or sending emails. Each intent is paired with a pre-defined Arabic example, ensuring that the assistant generates contextually relevant and linguistically consistent responses. The system uses prompt engineering to define the behavior of the assistant in four key steps:

1. understanding the user's intent.
2. determining whether the user is seeking information or an action.
3. executing the appropriate action based on the example
4. consistently responding in Arabic.

By using LangChain's LLMChain and ChatPromptTemplate, the assistant dynamically generates responses tailored to the user's input and maintains a coherent conversation history. This method allows for versatile interactions, as the assistant can switch between intents, maintain context from previous exchanges, and seamlessly handle both information requests and action-based commands.

6 Challenges and Solutions

6.1 Data Collection Challenges

Throughout the development of this system, several challenges were encountered, spanning data collection, model training, and integration. For **data collection**, two datasets were essential: one for the LSTM model, which required manually recording whistling noises to serve as the trigger word, and another for intent classification, which combined custom-written data and data generated using ChatGPT. The process of recording the audio for the trigger word was both time-consuming and prone to errors. Additionally, before settling on whistling, an Arabic word "اصحي" was considered, but the model's performance was subpar, likely due to the phonetic complexity of the word, leading to ambiguity and poor recognition. This issue highlights how certain words might be harder for models to recognize due to their similarity to common sounds or noise in the environment. In terms of solutions, selecting a distinctive non-speech sound like whistling solved the problem, as it reduced confusion with speech or ambient noise. The dataset for intent prediction was small, which could affect model performance. This limitation could be addressed by using data augmentation techniques to expand the dataset or by collecting more user data to improve diversity.

6.2 Model Training Challenges

For **model training**, the primary challenge was the resource intensity. The LSTM model, used for trigger-word detection, had to be trained from scratch locally, which is computationally expensive. Similarly, fine-tuning the pre-trained BERT model for intent classification was performed on Google Colab due to its high resource demands. A solution to this could be leveraging cloud-based GPUs or TPUs, or experimenting with lighter architectures that could reduce training time and resource requirements.

6.3 Integration Challenges

Finally, **integration challenges** emerged when merging the various components—trigger-word detection, speech-to-text, intent classification, and response generation—into a single functional system. The development of driver code to harmonize these modules was complex. One way to mitigate this would be using a modular development approach with clear interfaces between components, enabling easier debugging and scalability during integration. By addressing these challenges, the system's robustness and functionality were greatly improved.

7 Future Work and Improvements

7.1 Potential Enhancements

- **Enhanced Speech Recognition:** Explore alternative or complementary Arabic speech-to-text solutions that are specifically optimized for dialectal variations and

different accents within the Arabic language, which could improve the system's transcription accuracy.

- **Multilingual Support:** Expand the system's capabilities to handle multiple languages, enabling it to switch between Arabic and other languages based on the user's input, broadening the assistant's usability.

7.2 Additional Features

- **Action Execution:** Implement the ability for the smart assistant to perform actions based on user commands. For example:
 - **Home Automation:** The assistant could control smart home devices like lights, thermostats, and security systems upon request.
 - **Task Management:** The assistant could create, update, and manage to-do lists, calendar events, or reminders solely based on user commands.
 - **Information Retrieval:** The assistant could retrieve and provide information such as weather updates, news summaries, or directions.
 - **Hands-Free Calling:** Hands-free calling allows users to make calls to anyone on their contact list via voice commands.
 - **Service Integration:** The assistant could interact with external services such as booking a ride, ordering food, or sending messages on behalf of the user.

8 Conclusion

8.1 Summary of Achievements

The Arabic Smart Assistant project successfully developed a system that integrates several advanced machine learning and NLP technologies to create a functional and responsive assistant for Arabic-speaking users. Key achievements include:

- **Effective Trigger Word Detection:** A custom-trained LSTM model was successfully implemented to detect a whistling trigger word, enabling the assistant to respond promptly to user commands.
- **Accurate Arabic Speech Recognition:** Google Speech-to-Text was effectively utilized to transcribe spoken Arabic into text, providing a reliable foundation for further processing.
- **Precise Intent Classification:** A fine-tuned Arabic BERT model was employed to classify the transcribed text into specific intents, allowing the assistant to understand and respond appropriately to various user queries.
- **Contextually Relevant Responses:** A LLAMA large language model, combined with prompt engineering and the LangChain framework, was used to generate coherent and contextually relevant Arabic responses, enhancing the user experience.

- **Seamless User Interaction:** The assistant delivered the final output to the user through Google Text-to-Speech, ensuring a smooth and natural conversational flow.

8.2 Final Thoughts

The Arabic Smart Assistant project represents a significant step forward in the development of AI-driven language understanding and interaction for Arabic speakers. The assistant's ability to recognize, understand, and respond to Arabic speech opens new possibilities for personalized and context-aware user interactions. Future enhancements and additional features, such as action execution, will further elevate the assistant's utility and user engagement, making it an indispensable tool in daily life.

9 Appendices

9.1 Source Code

GitHub Repo: <https://github.com/adham137/Arabic-Voice-Assistant>

9.2 Colab Notebook

Used to train the BERT intent model:

<https://colab.research.google.com/drive/1ughLazoCppDFmUFDNuQcUcuHM0opGN4D>.